

Markov Decision Processes and Reinforcement Learning

Day 3 — Computing Optimal Value Functions and Policies

Luiz Felipe Martins

Department of Mathematics and Statistics
Cleveland State University

January 28, 2026

Outline

- 1 Review
- 2 Policy Iteration
- 3 Policy evaluation by Monte Carlo methods
- 4 Temporal Difference (TD) Learning
 - TD Prediction

Optimal Bellman Equation

- The **optimal Bellman equation** is:

$$V^*(s) = \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}} p(s' | s, a)[r(s, s', a) + \gamma V^*(s')]$$

- **Verification Theorem:** If V^* is a solution of the optimal Bellman equation, then we know it is the optimal value function
- In this case, an optimal deterministic policy can be obtained by **one-step lookahead**:

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}} p(s' | s, a)[r(s, s', a) + \gamma V^*(s')]$$

Value iteration algorithm

- Let $V_0(s)$ be initialized arbitrarily for $s \in \mathcal{S}$.
- Iterate:

$$V_{n+1}(s) = \max_{a \in \mathcal{A}(s)} \left\{ \sum_{s' \in \mathcal{S}} p(s, s' | a) [r(s, s', a) + \gamma V_n(s')] \right\},$$

except that we let $V_{n+1}(s) = 0$ if s is a terminal state.

Policy Iteration Algorithm

- Let π_0 be a randomized initial policy, with corresponding value function $V_0^\pi(\cdot)$.
- Iterate for $n = 0, 1, 2, \dots$:
 - Compute a new policy:

$$\pi_{n+1}(s) = \arg \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}} p(s' | s, a) [r(s, s', a) + \gamma V_n^\pi(s')]$$

- Compute the value function of π_{n+1} , $V^{\pi_{n+1}}(\cdot)$
- Stopping criteria:
 - $\pi_{n+1}(s) = \pi_n(s)$ for all $s \in \mathcal{S}$
 - $\max_{s \in \mathcal{S}} |V^{\pi_{n+1}}(s) - V^{\pi_n}(s)| < \epsilon_{\text{tol}}$
- A policy computed as above is called **greedy**.

Policy Iteration example

n	$\pi_n(0)$	$\pi_n(1)$	$V_n(0)$	$V_n(1)$
0	b	a	33.121	34.515
1	a	b	71.250	63.571
2	a	b	71.250	63.571

- In general, policy iteration converges much faster than value iteration.
- In fact, for a finite MDP, policy improvement will “converge” in a finite number of steps, since the total number of deterministic policies is finite.
- This assumes that:
 - Transition probabilities are known.
 - The whole state is observed.
 - Rewards are a discrete subset of \mathbb{R} .
 - $V_n(s)$ is computed exactly.

Approximate Policy Iteration

- Policy iteration starts with policies π_n that are far from optimal.
- It is wasteful to compute $V^{\pi_n}(s)$ exactly in the beginning.
- Instead, we compute only an *approximation* $\tilde{V}^{\pi_n}(s)$ in each step. For example, use a Jacobi or Gauss-Seidel iteration with a small number of relaxations.
- We don't even need a good approximation!
- Stopping condition:
 - Policy stabilizes: $\pi_{n+1}(s) = \pi_n(s)$ for all states s *and*
 - Value function estimates stabilize: $\tilde{V}^{\pi_{n+1}}(s) \approx \tilde{V}^{\pi_n}(s)$
- Compute a high-precision approximation of $V^{\pi_n}(s)$ for the final policy.
- Optionally, do one more policy improvement step to guarantee we have an optimal policy.

Approximate Policy Iteration Example

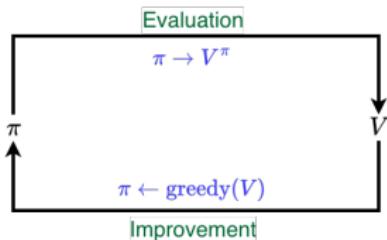
Approximate Policy Iteration for Tri-state MDP with 10 Gauss-Seidel relaxations for each policy improvement step:

n	$\pi_n(0)$	$\pi_n(1)$	$\tilde{V}_n(0)$	$\tilde{V}_n(1)$	$\max_s \tilde{V}_n(s) - \tilde{V}_{n-1}(s) $
0	b	a	32.59054893	34.02505034	—
1	a	b	70.18751040	62.82240404	3.760×10^1
2	a	b	71.22266853	63.55216067	1.035×10^0
3	a	b	71.24929693	63.57093292	2.663×10^{-2}
4	a	b	71.24998191	63.57141582	6.850×10^{-4}

High precision value function for final policy, obtained with 200 Gauss-Seidel relaxations:

$$V(0) = 71.25000000, \quad V(1) = 63.57142857$$

Generalized Policy Iteration



- Policy iteration alternates two stages:
 - **Evaluation:** Given a policy π , evaluate to V^π .
 - **Improvement:** Given a value function V , compute a policy π greedy with respect to V .
- **Generalized policy iteration (GPI)** algorithm: interleaves these two processes arbitrarily.
 - Value iteration: policy improvement step is performed after $V(s)$ is updated for each state.
 - Asynchronous DP: the two processes are carried out in parallel.

Why we need more general evaluation methods?

- From now on, we assume that we have a policy π (it may be deterministic or randomized), and we want to estimate its state value function $V(s)$.
- The methods seen so far are *model based*: the transition probabilities and rewards are known, and can't change.
- We want methods that can handle:
 - **Non-model based problems:** transitions and rewards are not known.
 - **Dynamic environments:** transitions and rewards can change with time.
- We need agents that can **learn** by interacting with the environment.

Monte Carlo methods

- **Monte Carlo** methods are used to analyze complex systems using computer-based simulations. They find applications in many areas of science.
- In many applications, we want to estimate the *expected value* of a random variable.
- We simulate a large number of realizations of the random variable, and use the average of the outcomes as an estimate for the expected value.
- Legend has that they are called “Monte Carlo methods” because of the famous casino in Monaco.

Monte Carlo Monaco



Monte Carlo Recife



Estimating the value function for a policy — Naive algorithm

- Input:
 - π , a policy to be evaluated.
- Initialization:
 - $V(s)$ initialized arbitrarily.
- For each state s :
 - Let `returns` = []
 - Let $S_0 = s$
 - Repeat a large enough number of times:
 - Generate an episode following π : $(S_0, R_0, A_0), \dots, (S_T, R_T, A_t)$
 - Let $G = \sum_{t=0}^T \gamma^t R_{t+1}$
 - Append G to `returns`
 - Let $V(s) = \text{average}(G)$

Example

- MDP has states $\{0, 1, 2, \dots, f\}$ and actions $\{a, b\}$. State f is terminal.
- The following are the results of three runs of the MDP ($\gamma = 1$):

Run 1: $(0, a, -), (4, b, 12), (2, a, 5), (f, -, 3)$ Total return: 20

Run 2: $(0, b, -), (0, b, 5), (f, -, 5)$ Total return: 10

Run 3: $(0, a, -), (3, a, 11), (1, a, 2), (f, -, 2)$ Total return: 15

- From this we get the estimate:

$$V(0) \approx \frac{1}{3}(20 + 10 + 15) = 15$$

- Repeat this procedure for all possible initial states.

Being more efficient

- To make things simpler, from now on we will consider only *episodic tasks*.
- Suppose we have a run of an episode of length n starting in the state S_0 :

$$(S_0, -), (S_1, R_1), (S_2, R_2), \dots, (S_n, R_n)$$

- From this we get one value of the return for the initial state S_0 , which is *one data point* that is used in the estimate of $V(S_0)$.
- However, we can consider a run that starts at time $t = 1$:

$$(S_1, R_1), (S_2, R_2), \dots, (S_n, R_n)$$

- This gives me a data point for estimating $V(S_1)$, if $S_1 \neq S_0$!
- Be careful adjusting the discount factor!

First-visit MC method

- Input:
 - π , a policy to be evaluated.
- Initialization.
 - $V(s)$, initialized arbitrarily.
 - `returns(s)`, an empty list for each state s .
- Repeat a large enough number of times:
 - Generate an episode following π : $(S_0, A_0, R_0), \dots, (S_T, A_T, R_T)$
 - Let $G = 0$
 - Loop for $t = T - 1, T - 2, \dots, 0$:
 - $G = \gamma G + R_{t+1}$
 - If S_t does not appear in S_0, S_1, \dots, S_{t-1} , append G to `returns(S_t)`
- Let $V(s) = \text{average}(\text{returns}(s))$ for $s \in \mathcal{S}$

First-visit MC method — Example

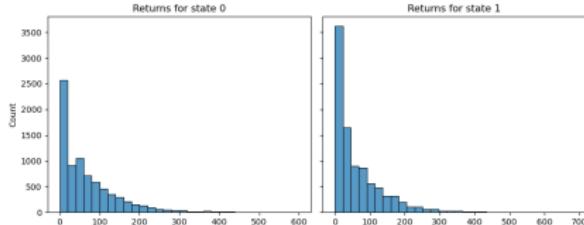
- For the Tri-state MDP, with policy $\pi(0) = a$, $\pi(1) = b$, simulating 10000 episodes we get the estimates:

$$V(0) = 71.18310, \quad V(1) = 63.25460$$

- Actual values:

$$V(0) = 71.25000, \quad V(1) = 63.57143$$

- Distribution of returns:



- $\text{std}[V(0)] = 74.77556, \text{ std}[V(1)] = 74.56148$

Every-visit MC method

- Input:
 - π , an admissible policy to be evaluated.
- Initialization.
 - $V(s)$, initialized arbitrarily.
 - $\text{returns}(s)$, an empty list for each state s .
- Repeat a large enough number of times:
 - Generate an episode following π : $(S_0, A_0, R_0), \dots, (S_T, A_T, R_T)$
 - Let $G = 0$
 - Loop for $t = T - 1, T - 2, \dots, 0$:
 - $G = \gamma G + R_{t+1}$
 - Append G to $\text{returns}(S_t)$
- Let $V(s) = \text{average}(\text{returns}(s))$ for $s \in \mathcal{S}$

Every-visit MC method — Example

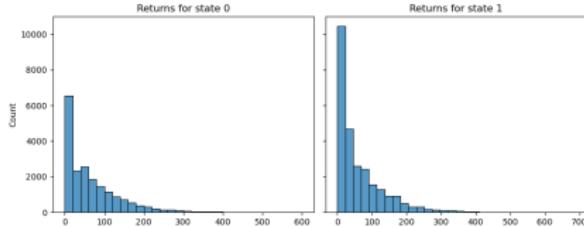
- For the Tri-state MDP, with policy $\pi(0) = a$, $\pi(1) = b$, simulating 10000 episodes we get the estimates:

$$V(0) = 69.84248, \quad V(1) = 62.18433$$

- Actual values:

$$V(0) = 71.25000, \quad V(1) = 63.57143$$

- Distribution of returns:



- $\text{std}[V(0)] = 70.86432, \text{ std}[V(1)] = 73.64984$

Dynamic update of averages

- Suppose that we have a simulation that gives us a sequence return values G_1, G_2, \dots, G_N . We use the average to estimate the expected return:

$$V_N = \frac{1}{N} \sum_{j=1}^N G_N$$

- Now, we compute a new value, G_{N+1} . How should we adjust the average? A simple computation shows that:

$$V_{N+1} = V_N + \frac{1}{N+1} [G_{N+1} - V_N]$$

- In control theory, this process is called *tracking*: The values of V_N are updated in a way that they “track” the values of G_N .
- Engineers like this because *the dynamics of the process G_N may change*, and the values of V_N can still try to “follow” it.

First-visit Monte Carlo with dynamic updates

- Input:
 - π , a policy to be evaluated.
- Initialization.
 - $V(s)$, initialized arbitrarily.
 - $\text{visits}(s) = 0$ for each state s .
- Repeat a large enough number of times:
 - Generate an episode following π : $(S_0, A_0, R_0), \dots, (S_T, A_T, R_T)$
 - Let $G = 0$
 - Loop for $t = T - 1, T - 2, \dots, 0$:
 - $G = \gamma G + R_{t+1}$
 - If S_t does not appear in S_0, S_1, \dots, S_{t-1} , increment $\text{visits}(s)$ by one and update:

$$V(s) \leftarrow V(s) + \frac{1}{\text{visits}(s)} [G - V(s)]$$

Target tracking

- The update rule:

$$V_{N+1} \leftarrow V_N + \frac{1}{N+1}[G_{N+1} - V_N]$$

is an example of *tracking* a non-stationary sequence G_N .

- The general form of this update rule is:

```
new_estimate = old_estimate + step_size[target - old_estimate]
```

- A wide variety of learning algorithms fit this pattern.

TD Learning, Monte Carlo and DP

- TD Learning is a combination of Dynamic Programming and Monte Carlo methods.
- Monte Carlo aspect: learn from experience, does not need a model for the environment dynamics.
- DP aspect: update estimates based on other learned estimates (bootstrap).
- TD is the basis for many current RL algorithms.

Framework for policy evaluation

- Suppose we have:
 - An estimate $V(s)$ for the state value function for a policy π .
 - An episode S_0, S_1, \dots, S_T generated with the policy π .
- Denote by G_t the total discounted return for this episode after time t :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-t-1} R_T$$

- For every visited state S_t , we update $V(S_t)$ to move in the direction of G_t .
- Different update strategies lead to distinct algorithms: Monte Carlo, TD(0), etc.

From Monte Carlo TD(0)

- Monte Carlo update rule:

$$V(S_t) \leftarrow V(S_t) + \frac{1}{\text{visits}(S_t)} [G_t - V(S_t)]$$

- Tracking with constant stepsize:

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

- TD(0): Instead of waiting for a whole episode to end, we use a one-step lookahead estimate of the value function:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

TD relationship to DP

- Bellman Equation for state value function V of a policy π :

$$V(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

The policy π may be stochastic or deterministic.

- Value iteration is a *fixed point iteration* for the exact value of $V(s)$:

$$V(s) \leftarrow \mathbb{E}_{\pi}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

- TD(0) is a *stochastic fixed point iteration* that tracks a one-step lookahead:

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

where S_0, S_1, \dots, S_T is an episode generated following policy π .

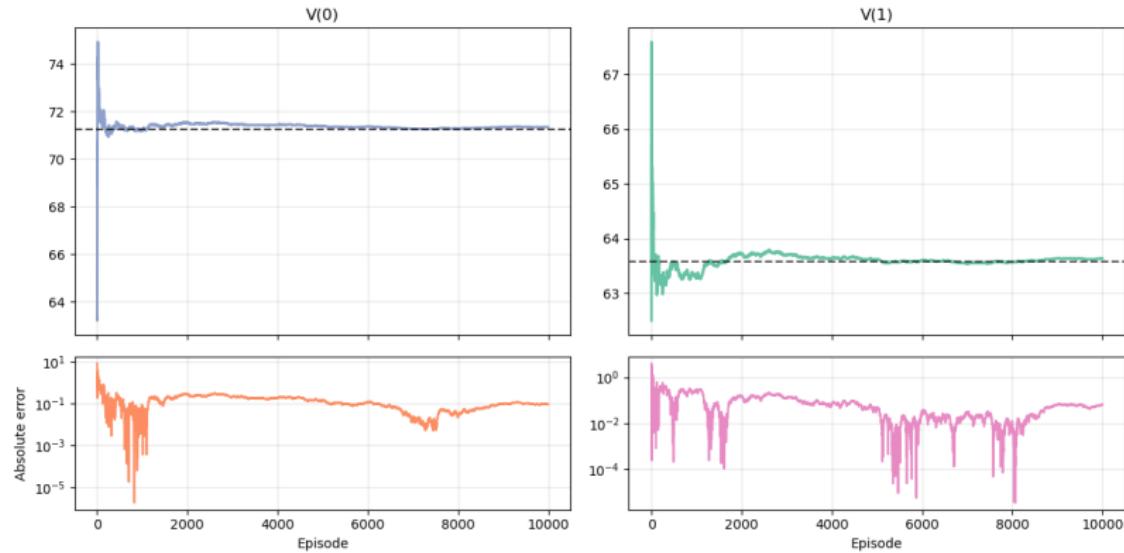
TD(0), Monte Carlo and DP

$$V_{\text{new}}(s) = V_{\text{old}}(S_t) + \alpha[r + \gamma V_{\text{old}}(s') - V_{\text{old}}(s)]$$

- Relationship to Monte Carlo:
 - Can learn directly from experience, independently on a model.
 - Monte Carlo algorithms have no bias and high variance.
 - TD(0) has bias controlled by the step size α and low variance.
- Relationship to DP
- TD methods update estimates based on already learned estimates (bootstrap)
- The TD(0) update rule is a form of *stochastic fixed point iteration*

Estimation by the Monte Carlo method

- Monte Carlo estimates for the Tri-state MDP with policy $\pi(0) = a$, $\pi(1) = b$ with 10000 episodes.
- Do 100 runs of the method and average.

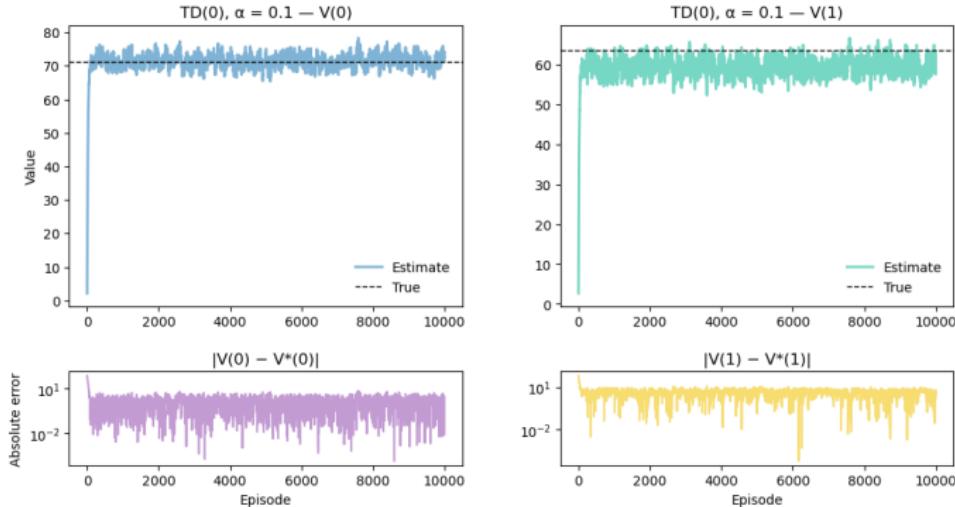


Interpretation of results for MC

- The value function estimates exhibit large transient oscillations, but stabilize around the true values. This behavior is consistent with an unbiased estimator.
- The absolute errors stabilize around 10^{-1} .
- The errors do not converge to zero at finite time due to the intrinsic variance of Monte Carlo returns and the finite number of effective samples per state.
- The magnitude of the error is governed by the standard deviation of the return distribution (empirically on the order of 70 in this model), with error scaling approximately as σ_G/\sqrt{N} .
- The sharp “dips” in the error correspond to episode indices where the mean estimate of $V(s)$ happens to be very close to the true value; this reflects bias cancellation rather than simultaneous accuracy across runs and is expected behavior.

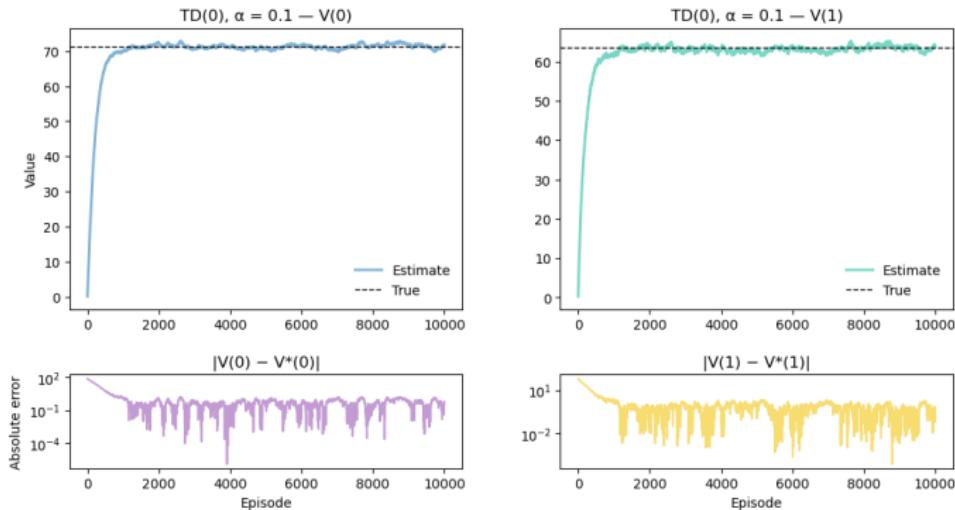
Estimation using TD(0) — $\alpha = 0.1$

- TD(0) estimates for the Tri-state MDP with policy $\pi(0) = a$, $\pi(1) = b$ with $\alpha = 0.1$ and 10000 episodes.
- Do 100 runs of the method and average.



Estimation using TD(0) — $\alpha = 0.01$

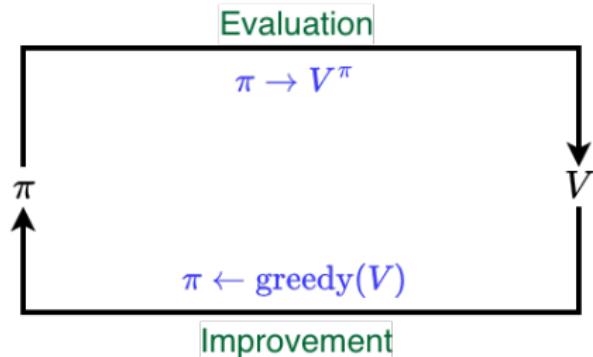
- TD(0) estimates for the Tri-state MDP with policy $\pi(0) = a$, $\pi(1) = b$ with $\alpha = 0.01$ and 10000 episodes.
- Do 100 runs of the method and average.



Interpretation of results for TD(0)

- TD(0) converges rapidly to a neighborhood of the true value function.
- With a fixed step-size, TD(0) does not converge pointwise, but to a stationary distribution around a biased fixed point.
- The stationary bias is *state-dependent* and need not be the same for all states.
- States with lower visitation frequency or closer to termination exhibit both higher variance and larger apparent bias.
- The variance of the estimator does not vanish and is controlled by the step-size: larger step-sizes yield faster learning but higher steady-state variance.

From estimation to control



We have not explored the “Evaluation” branch of the cycle. Next time we will see how this can be incorporated in a policy optimization algorithm.